


Les index PostgreSQL

ET LEUR INTÉGRATION AVEC DJANGO

FRANÇOIS FREITAG

LES EMPLOIS DE L'INCLUSION
GIP INCLUSION

26 AOÛT 2024



- 1 Le scan séquentiel
- 2 Indexer les données
- 3 Les types d'index
- 4 Méthodes d'indexation
- 5 Django

LES DONNÉES, LA REQUÊTE

```
CREATE TABLE test1 (  
    id INTEGER,  
    content VARCHAR  
);
```

```
SELECT content FROM test1 WHERE id = constant;
```

EXPLAIN ANALYZE

ANALYZE : Carry out the command and show actual run times and other statistics.

```
EXPLAIN ANALYZE SELECT content FROM test1 WHERE id = 12;
```

QUERY PLAN

```
Seq Scan on test1 (cost=0.00..25.88 rows=6 width=32)
      (actual time=0.016..0.017 rows=0 loops=1)
```

```
    Filter: (id = 12)
```

```
Planning Time: 2.165 ms
```

```
Execution Time: 0.098 ms
```

```
(4 rows)
```

- 1 Le scan séquentiel
- 2 Indexer les données**
- 3 Les types d'index
- 4 Méthodes d'indexation
- 5 Django

INDEXER LES DONNÉES

De la même manière qu'un livre contient souvent un index trié par ordre alphabétique, la base de données crée un index trié des données présentes dans une table.

```
CREATE INDEX test1_id_index ON test1 (id);
```

```
EXPLAIN ANALYZE SELECT content FROM test1 WHERE id = 12;
```

QUERY PLAN

```
-----  
Bitmap Heap Scan on test1 (cost=4.20..13.67 rows=6 width=32)  
    (actual time=0.010..0.011 rows=0 loops=1)  
    Recheck Cond: (id = 12)  
    -> Bitmap Index Scan on test1_id_index  
        (cost=0.00..4.20 rows=6 width=0)  
        (actual time=0.008..0.008 rows=0 loops=1)  
        Index Cond: (id = 12)  
Planning Time: 0.118 ms  
Execution Time: 0.053 ms  
(6 rows)
```

- Espace disque supplémentaire
- Maintenance de l'index lors des opérations en écriture (INSERT, UPDATE, DELETE)

- 1 Le scan séquentiel
- 2 Indexer les données
- 3 Les types d'index**
- 4 Méthodes d'indexation
- 5 Django

B-TREE

- Index par défaut de PostgreSQL
- Equality (=), range queries (<, >, BETWEEN, IN), NULL

root node

non-leaf nodes

leaf nodes

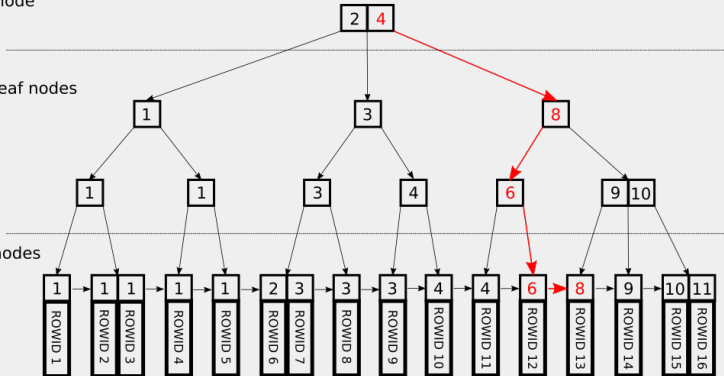


Figure – source :

<https://www.qwertee.io/blog/postgresql-b-tree-index-explained-part-1/>

HASH

- Hash indexes store a 32-bit hash code derived from the value of the indexed column.
- Equality (=), not inequality.
- More compact and can be faster than a B-Tree index

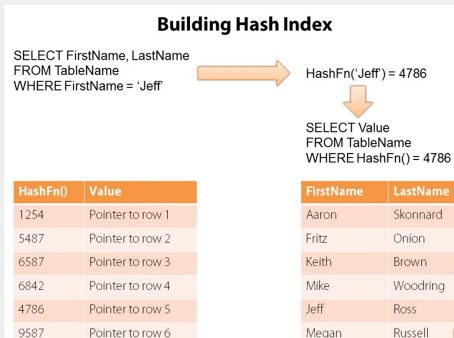


Figure – source :

<https://gss-portal.com/knowledgebase/129/Hash-Index-Part-4.html/>

GIST - GENERALIZED SEARCH TREE

- A balanced, tree-structured access method, that acts as a base template in which to implement arbitrary indexing schemes.
- Works on ranges (containment <@, overlap &&, strictly left/right «, », etc)
- Handles
 - ▶ ranges (integer, dates, IP subnets, etc),
 - ▶ geographical data,
 - ▶ string distance,
 - ▶ bioinformatics,
 - ▶ nearest neighbors,
 - ▶ and generally what domain experts can modelize with a set of standard operations.

SP-GiST - SPACE-PARTITIONED GiST

- Same as GiST, but allows for unbalanced structures.
- Can be much faster for domains where partitioning makes sense.

- Appropriate for data values that contain multiple component values, such as arrays or JSON.
- Equality (=), containment (<@, @>), overlap (&&).

⇒ C'est ce que les `emails_email` utilisent, pour rechercher une adresse email dans la liste des champs `To`, `Cc` et `Bcc`.

BRIN - BLOCK RANGE INDEX

- Designed for handling very large tables in which certain columns have some natural correlation with their physical location within the table.
- e.g. storing a store's sale orders might have a date column on which each order was placed
- Lossy, must recheck the tuples and discarding those that do not match the query conditions
- Very small, avoids scanning large parts of the table

- A Bloom filter is a space-efficient data structure that is used to test whether an element is a member of a set.
- Lossy, takes a signature based on the data and returns tuples matching the signature
- Many attributes and queries test arbitrary combinations of them
- Equality (=), not inequality.
- Requires an extension
- *Example on PostgreSQL documentation*

- 1 Le scan séquentiel
- 2 Indexer les données
- 3 Les types d'index
- 4 Méthodes d'indexation**
- 5 Django

FUNCTIONAL INDEX

```
CREATE INDEX test1_upper_content_index ON test1 (UPPER(content));
```

```
EXPLAIN ANALYZE SELECT content FROM test1 WHERE UPPER(content)='FOO';
```

QUERY PLAN

```
Bitmap Heap Scan on test1 (cost=4.20..13.68 rows=6 width=32)
    (actual time=0.010..0.011 rows=0 loops=1)
    Recheck Cond: (upper((content)::text) = 'FOO'::text)
    -> Bitmap Index Scan on test1_upper_content_index
        (cost=0.00..4.20 rows=6 width=0)
        (actual time=0.008..0.008 rows=0 loops=1)
        Index Cond: (upper((content)::text) = 'FOO'::text)
Planning Time: 0.142 ms
Execution Time: 0.055 ms
```

⚠ Requires a functional query.

```
EXPLAIN ANALYZE SELECT content FROM test1 WHERE content='FOO';
```

QUERY PLAN

```
-----  
Seq Scan on test1  (cost=0.00..25.88 rows=6 width=32)  
      (actual time=0.017..0.018 rows=0 loops=1)  
    Filter: ((content)::text = 'FOO'::text)  
Planning Time: 0.865 ms  
Execution Time: 0.105 ms  
(4 rows)
```

Module `citext` pour les adresses e-mail des emplois de linclusion.

- Comparaison non sensible à la casse
- Pas besoin d'appeler de fonction sur les valeurs
- Fonctionne avec `UNIQUE`
- Les emails n'ont pas besoin d'une normalisation de l'unicode, car les emplois ne supportent pas l'*Email Address Internationalization*

Refs ce commit

PARTIAL INDEX

```
CREATE INDEX test1_large_ids_index ON test1 (id) WHERE id >= 300000;
```

```
EXPLAIN ANALYZE SELECT * FROM test1 WHERE id > 30000;
```

QUERY PLAN

```
Bitmap Heap Scan on test1 (cost=7.43..22.72 rows=423 width=36)
    (actual time=0.002..0.003 rows=0 loops=1)
    Recheck Cond: (id > 3000000)
    -> Bitmap Index Scan on test1_composite_index
        (cost=0.00..7.32 rows=423 width=0)
        (actual time=0.001..0.002 rows=0 loops=1)
        Index Cond: (id > 30000)
Planning Time: 0.047 ms
Execution Time: 0.016 ms
(6 rows)
```

COMPOSITE INDEX

```
CREATE INDEX test1_composite_index ON test1 (id, content);
```

```
EXPLAIN ANALYZE SELECT * FROM test1 WHERE id > 30000;
```

QUERY PLAN

```
Bitmap Heap Scan on test1 (cost=7.43..22.72 rows=423 width=36)
    (actual time=0.002..0.003 rows=0 loops=1)
    Recheck Cond: (id > 3000000)
    -> Bitmap Index Scan on test1_composite_index
        (cost=0.00..7.32 rows=423 width=0)
        (actual time=0.001..0.002 rows=0 loops=1)
        Index Cond: (id > 30000)
Planning Time: 0.047 ms
Execution Time: 0.016 ms
(6 rows)
```

MOST SPECIALIZED INDEX

```
EXPLAIN ANALYZE SELECT * FROM test1 WHERE id > 300000;
```

QUERY PLAN

```
Bitmap Heap Scan on test1 (cost=4.24..19.53 rows=423 width=36)
    (actual time=0.011..0.012 rows=0 loops=1)
    Recheck Cond: (id > 3000000)
    -> Bitmap Index Scan on test1_large_ids_index
        (cost=0.00..4.13 rows=423 width=0)
        (actual time=0.008..0.008 rows=0 loops=1)
        Index Cond: (id > 3000000)
Planning Time: 0.178 ms
Execution Time: 0.047 ms
(6 rows)
```

- 1 Le scan séquentiel
- 2 Indexer les données
- 3 Les types d'index
- 4 Méthodes d'indexation
- 5 Django**

```
class Index(*expressions, fields=(), name=None, db_tablespace=None,  
            opclasses=(), condition=None, include=None)
```

```
# Un index fonctionnel ?
```



```
class Index(*expressions, fields=(), name=None, db_tablespace=None,
            opclasses=(), condition=None, include=None)

# Un index fonctionnel ?
Index(Lower("title").desc(), name="lower_title_idx")
# Un index composite ?
```

```
class Index(*expressions, fields=(), name=None, db_tablespace=None,  
            opclasses=(), condition=None, include=None)
```

```
# Un index fonctionnel ?
```

```
Index(Lower("title").desc(), name="lower_title_idx")
```

```
# Un index composite ?
```

```
Index("title", "pub_date", name="title_date_idx")
```

```
# Un index partiel ?
```

```
class Index(*expressions, fields=(), name=None, db_tablespace=None,
            opclasses=(), condition=None, include=None)

# Un index fonctionnel ?
Index(Lower("title").desc(), name="lower_title_idx")
# Un index composite ?
Index("title", "pub_date", name="title_date_idx")
# Un index partiel ?
Index(
    "pub_date",
    condition=Q(pub_date__lt=datetime.date(2020, 1, 1)),
    name="published_before_2020_idx",
)
```

```
# django.contrib.postgres.indexes:
class BloomIndex(*expressions, length=None, columns=(), **opts): ...
class BrinIndex(*expressions, autosummarize=None,
                pages_per_range=None, **opts): ...
class BTreeIndex(*expressions, fillfactor=None, **opts): ...
class GinIndex(*expressions, fastupdate=None,
               gin_pending_list_limit=None, **opts): ...
class GistIndex(*expressions, buffering=None,
               fillfactor=None, **opts): ...
class HashIndex(*expressions, fillfactor=None, **opts): ...
class SpGistIndex(*expressions, fillfactor=None, **opts): ...
```

MERCI DE VOTRE ATTENTION

AVEZ-VOUS DES QUESTIONS ?